

THE PROBLEM

Oncologists treating complex, heavily pretreated patients need to match a specific clinical picture against the literature and open trials — fast, and again every time that patient's status changes. Three things get in the way: the relevant evidence is scattered across a paper's sections (eligibility, methods, results, subgroup analyses), not sitting in one paragraph; a real patient profile has eight or more clinical variables that get averaged into noise by a single search query; and the language on the chart ("no longer a surgical candidate") rarely matches the language in the literature ("unresectable"). Add hundreds of new oncology papers published every week, and no clinician can screen all of it manually — once, let alone continuously — on top of a full clinic load.

WHAT PAXIS DOES

Paxis ingests oncology literature into a structured, section-aware evidence base and lets a clinician submit a patient — as a structured record or in plain language — and get back matched, ranked studies with the specific evidence extracted, not just a list of links. The patient isn't a one-off query: their diagnosis, biomarkers, and treatment response stay on record, so a clinician can keep asking what's working and what isn't for that specific case over time — and Paxis keeps watching the literature in the background, surfacing a new matching study the moment it's published.

WHAT IT DOES THAT OTHERS CAN'T

Remembers the patient, not just the question

Diagnosis, biomarkers, and treatment response stay attached to the case. Ask what's working, what isn't, and why — Paxis answers in the context of that specific patient, not a generic query.

Never stops watching

Turn on alerts for a saved case and Paxis keeps checking new literature in the background, notifying you the moment a matching study is published — no manual re-search, ever.

Reads clinical language like literature does

An ontology-inference layer translates implicit chart language into the terms the literature actually uses, automatically — no manual translation, no missed matches.

Searches the whole patient, not one query

The profile is decomposed into individual clinical axes and searched in parallel across a vector index, structured clinical records, and patient-outcome frames — then merged.

Filters before it shows you anything

A dedicated relevance model re-ranks every candidate against the full patient profile, rejecting superficial keyword matches before they reach the clinician.

Covers the whole workflow, not just search

Intake, ongoing conversation, matching, treatment comparison, trial finder, and analytics — one product built around the patient, not a search box.

BUILT AND RUNNING

- **Live, deployed product** — not a prototype — running on dedicated cloud infrastructure
- **Real oncology literature corpus** already ingested and structured
- **Persistent per-patient records** with treatment-response tracking, feeding both conversation and matching
- **Standing case alerts** running in production — continuous monitoring, not one-time search
- **Multi-source retrieval pipeline** — vector search, structured clinical matching, and patient-outcome frames — operating together
- **Full clinical workflow shipped**: intake, matching, treatment comparison, trial finder, and analytics